

Let's Take a (Random) Walk

Soumya Ray

2019-02-24

Convergence of Binomial Probabilities¹

WE ALL INTUITIVELY UNDERSTAND that when a fair coin (50% heads, 50% tails) is flipped many times, the number of heads versus tails should start to converge to being equal. For example, let's say we get one point (+1) for every head but lose one point (-1) for every tail. If we flip a fair coin 500 times, then which each flip we expect the *average* (mean) of all flips thus far to converge towards zero. Let's simulate this and plot the results in R.

```
n <- 5000
flips <- 2*rbinom(n, 1, 0.5)-1 # a fair coin follows a simple binomial distribution
means <- c(seq(1:n), NA)
for (i in 1:n) { means[i] <- mean(flips[1:i]) }

plot(means, type="l", ylim=c(-1,1))
abline(h=0, col="red")
```

Just as we expected, we see that the average number of points quickly converges to 0 (zero). This is our common intuition about how probabilities average out over time.

INSTEAD OF AVERAGE PROBABILITIES, what if we considered total points? Let's play the game again: we gain one point for heads and lose one point for tails. How many *total points* (sum) would we make over time? Will it converge?

```
n <- 5000
flips <- 2*rbinom(n, 1, 0.5)-1
totals <- c(0)
for (i in 2:n) { totals[i] <- totals[i-1] + flips[i] }

plot(totals, type="l", ylim=c(-300, 300))
abline(h=0, col="red")
```

Rather than converging, it seems to just walk away! *Random walks* are one way of modeling many everyday phenomena that are affected by so many variables that they appear random in their progress: stock markets, weather fluctuations, ant foraging, brownian motion, and more².

¹ These examples are based on:
<http://nvenkataraman1.github.io/random-walks/>




Figure 1: Coin flipping probabilities

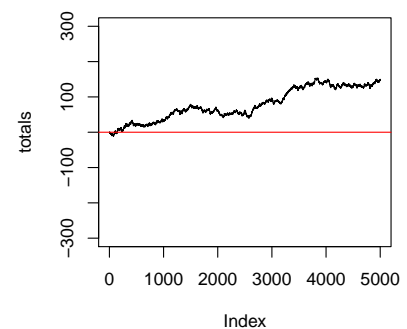


Figure 2: A random (binomial) walk


² See: http://en.wikipedia.org/wiki/Random_walk

BUT HOW FAR FROM THEIR STARTING POINTS do random walks go *on average*? Put another way, if many people played the 500-flips game we've described, would they find their totals points walking away in the same manner? It isn't easy to tell, so let's simulate many random walks and find out.

```
# Random binomial walk function
random_binomial_walk <- function(steps) {
  flips <- 2*rbinom(n, 1, 0.5)-1
  totals <- c(0)
  for (i in 2:steps) { totals[i] <- totals[i-1] + flips[i] }
  return(totals)
}

# Simulate six random walks of 5000 steps each
walks <- data.frame(matrix(0, nrow = 5000, ncol = 6))
for (i in 1:6) {
  walks[i] <- random_binomial_walk(5000);
}

# Plot out our six random walks
par(mfrow=c(2,3)) # setup plotting to 2rows x 3 columns
for (i in 1:6) {
  plot(walks[[i]], type="l", ylim=c(-200, 200))
  abline(h=0, col="red")
}
```



We first create a 'function' that creates binomial walks for us. It reuses the code we saw earlier. This function takes the number of 'steps' to walk as a parameter.

We first create a 'data.frame' that is the size of a 2x3 matrix. The we use a loop to run our random binomial walk function six times. The result of our six function calls are stored in the data.frame

We then plot the six walks. The 'mfrow()' function lets us put plots side-by-side in a grid pattern. At the end, we have to be careful to call 'mfrow()' again to return the graphics setting to plot 1x1 figures at a time.

Figure 3: Six Random Binomial Walks

```
par(mfrow=c(1,1)) # return graphics settings to 1 x 1
```

Even with a small number of random walks, we see they can go anywhere. It will be hard to plot a larger number of walks. So let's simulate a large number of walks and see a distribution of distances-from-zero where all the walks ended up. We will use 3000 walks, each having 2000 steps.

```
m <- 3000
n <- 2000
walks <- data.frame(matrix(0, nrow = n, ncol = m))
for (i in 1:m) { walks[i] = random_binomial_walk(n) }
final.points <- c(0)
for (i in 1:m) { final.points[i] <- walks[n,i] }
```

We can now plot the distribution of how many points away from zero (positive or negative) all the walks got to.

```
hist(final.points, breaks=20, prob=TRUE)
lines(density(final.points), lwd=3, col="blue")
```

Does this bell-shaped symmetric distribution look familiar? It is our old friend the *normal distribution* again!




Figure 4: Distribution of binomial walks